

YOGITA JAIN

AI & Data Engineer

yogitajain2510@gmail.com | +91 9165317750 | Pune, Maharashtra, India | [linkedin.com/in/yogita-jain](https://www.linkedin.com/in/yogita-jain)

PROFESSIONAL SUMMARY

Senior Data & AI Engineer with about 5 years of experience delivering scalable ETL pipelines, cloud data platforms, and enterprise analytics using AWS, Databricks, Palantir, PySpark, and SQL. Demonstrated business impact through \$600K infrastructure savings, \$1M compliance cost reduction, and 45% user-efficiency gains. Production-grade expertise in LLM applications, RAG systems, and AI agents on AWS Bedrock. Seeking Senior Data / AI Engineer roles in enterprise tech, fintech, or AI-first product companies.

TECHNICAL SKILLS

Languages: Python, SQL, PySpark

Data Engineering: ETL/ELT Pipelines, Data Modeling, Data Warehousing, Data Lakes, SCD Type 2, Data Integration, Batch & Streaming Pipelines, Delta Lake

Big Data & Platforms: Apache Spark, Databricks, Palantir Foundry

Cloud & Infrastructure: AWS (S3, Glue, Lambda, Redshift, Athena, Step Functions, EC2), Terraform

Generative AI & ML: Agentic AI, AWS Bedrock, Claude, LLMs, RAG Architecture, Prompt Engineering, LangChain, LangGraph, Embeddings, Vector Search, AI Agents, MLflow

Data Governance & Quality: Data Governance, Data Quality Frameworks, Metadata Management, Risk & Compliance, Data Mesh principles

Tools & Visualization: DBT, Git, GitHub Actions, Jira, Power BI, Jupyter Notebooks, VS Code

PROFESSIONAL EXPERIENCE

AI & Data Engineer | John Deere

Jun 2025 – Present

Product Delivery Data & Analytics

Tech Stack: PySpark, Python, SQL, Databricks, Palantir, AWS, Terraform, OpenAI API, AWS Bedrock, Claude, LLMs, RAG, Delta Lake

- Independently built a production-grade Generative AI agent using AWS Bedrock, Claude, OpenAI GPT, React, and RAG architecture — enabling semantic search and automated PLM query resolution, reducing mean resolution time by 60% and improving user efficiency by 45%.
- Architected and optimized scalable ETL pipelines using PySpark, SQL, Databricks, and Palantir for manufacturing and PLM analytics, processing 100 TB+ of enterprise data daily with high reliability and sub-hour SLA compliance.
- Orchestrated Palantir pipeline optimizations using linter-driven recommendations and automated scheduling, delivering approximately \$600K in annual infrastructure cost savings.
- Designed a centralized API-based automation framework for enterprise data product registration across governance and quality dimensions, eliminating ~200 hours/month of manual compliance effort.
- Resolved critical AWS cloud vulnerabilities and owned enterprise-wide security and risk compliance across all production environments, achieving zero critical findings in quarterly audits.
- Learned Terraform and independently executed production EC2 infrastructure migration across AWS accounts, contributing to full AWS account decommissioning and additional cost optimization.
- Deployed an Incomplete Parts Prediction ML solution at an enterprise hackathon for the PartsVerse team using MLflow-tracked predictive analytics to identify manufacturing gaps before downstream impact.

Cloud Data Engineer | Deloitte USI

Sep 2021 – May 2025

Lincoln Financial – Product Strategy

Tech Stack: PySpark, Python, SQL, AWS S3, Glue, Lambda, Amazon Redshift, Step Functions

- Designed and deployed batch and streaming data pipelines using AWS Glue, Lambda, and Step Functions for financial investment product analytics, reducing data latency by 40%.
- Automated sprint documentation and requirements workflows using Python scripts, speeding up the project timeline by 25% and freeing ~15 hours/week of business analyst effort.
- Spearheaded sprint planning and retrospectives for data migration modules, ensuring 99.9% data quality SLA across all pipelines.

NYL EDM – Cloud Migration

Tech Stack: PySpark, Python, SQL, AWS S3, Glue, Lambda, EC2, Athena, Step Functions, Amazon Redshift, Docker

- Reduced \$1M in business compliance costs by delivering the high-impact CCPA detokenization module within tight timelines, processing 50M+ records with zero data loss.
- Delivered 25+ optimization features, 120+ tables, and 10+ enhancements across ~25 sprints while managing 30+ interdependent work streams.
- Developed scalable, cost-effective ETL frameworks for migration of 250+ data sources; automated metadata configuration to boost ingestion efficiency by 90%, cutting onboarding time from days to hours.

NYL Service TZ Migration

Tech Stack: Amazon Redshift, DBT, SQL, EC2

- Converted on-prem HQL to Amazon Redshift SQL, boosting data storage and processing performance by 30–40%, and reducing monthly compute costs by ~\$80K.
- Transformed 50+ tables across 10+ sources using DBT models; built ELT workflows and optimized SQL queries, stored procedures, and scripts, cutting average query runtime by 35%.

DirectTV – Data Strategy

Tech Stack: SQL, Excel, MS Visio, MS Office

- Led Data Design work stream for 5+ high-priority business use cases supporting 10M+ subscriber records; designed Data Lineage diagrams and Source-to-Target mappings adopted as the team standard.
- Delivered end-to-end data documentation for 3 major reporting modules, reducing onboarding time for new analysts by approximately 30%.

Software Development Engineer Trainee | CRIS – Indian Railways

Jun 2020 – Nov 2020

DLW – Indian Railways

Tech Stack: Python, Django, SQL, PostgreSQL, Visual Studio Code

- Developed 15+ interactive web forms using Python, Django, and PostgreSQL for the DLW wing of Indian Railways, digitizing manual workflows for 500+ daily users.
- Led requirements gathering, blueprint design, and end-to-end testing for peer modules, reducing post-launch defect rate by 25%.

CERTIFICATIONS

- AWS Certified Generative AI Developer Professional
- AWS Certified Machine Learning Engineer Associate
- AWS Certified Data Engineer Associate
- AWS Certified Developer Associate
- AWS Certified AI Practitioner
- AWS Certified Cloud Practitioner

EDUCATION

Bachelors of Technology – Computer Science | Banasthali Vidyapith CGPA: 8.45 / 10

Jul 2017 – Jun 2021

AWARDS & RECOGNITION

- | | |
|--|----------|
| • Commendable Contribution Individual Award — John Deere | May 2026 |
| • WOW Performance Award — John Deere | Jan 2026 |
| • Data Engineering Performance Award — Deloitte USI | Mar 2025 |
| • AWS Community Initiative Award — Deloitte USI | Jan 2025 |
| • CCPA Development Applause Award — Deloitte USI | Jan 2024 |
| • Data Pipeline Performance Award — Deloitte USI | Jul 2023 |
| • Data Ingestion Performance Award — Deloitte USI | Dec 2022 |
| • Data Modelling Performance Award — Deloitte USI | Mar 2022 |